

Methodology

AI Career Compass | transparent rules for public US AI postings across experience levels

Snapshot	2026-09-21
Postings / companies	852 / 107
Coverage checked	198

Purpose

AI Career Compass is a dated, source-backed guide to public US AI postings. Its default flow asks about background, role interests, project skills, and preferences, then shows evidence-based statuses and next steps. It does not predict hiring success or estimate the labor market.

Data and scope

The snapshot contains 852 deduplicated public postings from 107 companies and 198 source checks. Sources: Ashby (309), Company careers (17), Greenhouse (522), Lever (4). Pathway labels at run time: early-career (9), explicit (42), unclear (801). Fresh official sources cover US technical AI roles at all experience levels; collection has no seniority or year ceiling.

Inclusion rules

Curated records are selected for full-time status, a US location or explicit Remote-US scope, core AI or ML responsibilities, and an active official application route. Ashby discovery checks FullTime; Greenhouse discovery does not establish employment terms. Both remain separately labeled unclear until qualification review. Pathways are explicit new-grad, zero-experience, early-career (0-2 industry years), or unclear (graduate eligibility not established). Title-based experience level is independent of graduate eligibility. Numerical year references are not validated minimums; unknown is never zero.

Extraction and review

Collectors retrieve official ATS or company-career pages and normalize IDs, URLs, locations, role family, qualifications, skills, salary, and sponsorship statements. Curated records receive source-specific review; broad discovery is automated and is not presented as individually reviewed. AI assists extraction and summaries. This is agent review, not human verification. Evidence paraphrases stay linked to the official page and missing values remain unknown.

Product logic

The four-step guide covers background, role interests, project skills, and preferences. Aligned, confirm, and gap statuses use only recorded conflicts or unknowns in stated degree, graduation, and experience conditions. Skills and preferences explain preparation and ordering but do not reject a person or create success odds. The beginner preset keeps explicit graduate evidence, entry labels, and unspecified titles without a year reference above two; users can switch levels. What-if views explore overlap. Remember on this device is an opt-in local browser profile; no applicant notes are imported.

Deduplication and analysis

Company plus source ID is the primary key, followed by canonical URL when needed. Role counts are postings, not headcount. Location totals count each posting once per listed state and keep Remote - US separate; states can overlap because one posting can list several states. Salary figures are disclosed base ranges only. Missing data is not a negative answer.

Limits and provenance

The earlier career/26fall dataset was used only for discovery and is biased toward Bay Area LLM roles; no applicant notes were imported. Fresh official sources broaden US coverage, but this remains a purposive public-posting sample rather than a census. ATS access, company selection, posting turnover, and wording differences create coverage bias, and a posting can close after the snapshot.

Reproduction

Run `generate_documents.py` against the frozen `dist/data.json` snapshot after the expanded partitions are merged. All counts are computed at run time and the script writes the methodology, reflection, demo script, data dictionary, and keyboard-driven presentation under `research/artifacts`.