

Methodology

AI Career Compass | visualizer rules and evidence boundaries

Snapshot	2026-09-21
Postings / companies	852 / 107
Source checks	198

Purpose and views

AI Career Compass supports Opportunity landscape for the filtered sample, Skill lab for local skill-mention exploration, and Job postings for the full posting reference list. The reference list exposes experience, skills, direction, location, and source without an application channel.

Canonical snapshot

The snapshot contains 852 unique postings from 107 companies and 198 company/source checks. Sources are Ashby (309), Company careers (17), Greenhouse (522), Lever (4). Counts are postings, not headcount; source links, review labels, and data-quality notes remain available in the Project files downloads.

Shared filters and measures

All three views share primary role direction, experience level, US region/state/remote, and graduate evidence. Job postings adds posting search and a canonical skill-mention filter that can persist into the charts and Skill lab. Landscape reports distributions and a skill-demand matrix. Skill lab reports the count and fraction of filtered postings meeting a per-posting skill target (fixed at 70%) by role, then previews before/after coverage when one skill is tried.

Experience labels

Experience level is one of four broad descriptive exploration buckets and is independent from graduate eligibility. Source evidence maps 0–2 years to entry; experienced individual contributors with 3+ years, supported independent scope, or a prior production track record to senior; explicit senior technical scope to staff; and people-management scope to manager. Ordinary experience familiarity alone is insufficient. Unsupported cases stay unspecified. The basis and explanation are recorded, while graduate flags and qualification paths are unchanged.

Skill interpretation

A posting meets the target when $\text{matched skills} \geq \text{ceil}(\text{distinct tracked skills} * \text{target})$. Fixed target: 70%; four of five tracked skills. Postings without skills are excluded from coverage denominators. Recommendations rank unselected skills by additional postings reaching the target, then normalized progress on unmet postings, then company breadth. Mention coverage is not proficiency, eligibility, readiness, hiring probability, or absence of an unmentioned skill. Automated extraction uses a fixed vocabulary. Sparse skill lists can inflate coverage; required, preferred, contextual and alternative mentions are weighted equally.

Evidence and privacy

Records retain official URLs, source IDs, timestamps, extraction notes, and review level. The sample is non-representative and state/skill counts can overlap. Selected skills stay in page memory and optional device-local storage; they are not sent to a server. Public CSV downloads retain the evidence fields.