

Methodology

AI Career Compass | visualizer rules and evidence boundaries

Snapshot	2026-09-21
Postings / companies	852 / 107
Source checks	198

Purpose and views

AI Career Compass supports Opportunity landscape for the filtered sample, Skill lab for local skill-mention exploration, and Job references for the full posting reference list. The reference list exposes experience, skills, direction, location, and source without an application channel.

Canonical snapshot

The snapshot contains 852 unique postings from 107 companies and 198 company/source checks. Sources are Ashby (309), Company careers (17), Greenhouse (522), Lever (4). Counts are postings, not headcount; source links, review labels, and data-quality notes remain available in the Project files downloads.

Shared filters and measures

All three views share primary role direction, experience level, US region/state/remote, and graduate evidence. Job references adds posting search and a canonical skill-mention filter that can persist into the charts and Skill lab. Landscape reports distributions and a skill-demand matrix. Skill lab reports the count and fraction of filtered postings mentioning at least one selected skill by role, then previews before/after coverage when one skill is tried.

Experience labels

Experience level is one of four broad descriptive exploration buckets and is independent from graduate eligibility. Source evidence maps 0–2 years to entry; experienced individual contributors with 3+ years, supported independent scope, or a prior production track record to senior; explicit senior technical scope to staff; and people-management scope to manager. Ordinary experience familiarity alone is insufficient. Unsupported cases stay unspecified. The basis and explanation are recorded, while graduate flags and qualification paths are unchanged.

Skill interpretation

A posting counts once when it mentions at least one selected skill. Recommendations exclude selected skills and rank by posting mentions, then company breadth. Mention coverage is not proficiency, eligibility, readiness, hiring probability, or absence of an unmentioned skill. Automated extraction uses a fixed vocabulary.

Evidence and privacy

Records retain official URLs, source IDs, timestamps, extraction notes, and review level. The sample is non-representative and state/skill counts can overlap. Selected skills stay in page memory and optional device-local storage; they are not sent to a server. Public CSV downloads retain the evidence fields.