

Reflection and limitations

AI Career Compass | what the evidence can and cannot say

Snapshot	2026-09-21
Postings / companies	852 / 107
Coverage checked	198

What the snapshot supports

The snapshot supports a reproducible description of 852 included postings across 107 companies, including role family, stated US locations, qualification pathways, selected skills, disclosed base salary, and public application links. Recorded pathway categories are dynamic (early-career: 9, explicit: 42, unclear: 801).

What the product supports

The four-step guide helps a beginner state a background, choose role interests, add project skills, and set preferences. Results explain aligned, confirm, or gap statuses with source evidence. A what-if view can show skill overlap and a nationwide search. These are evidence and preparation aids, not ML predictions, success odds, or recruiter decisions.

AI assistance and agent review

AI was used as an extraction aid for normalization, grouping, and short summaries. Curated records received source-specific checks for URL, title, US location, employment type where exposed, AI evidence, pathway, and application route; broad automated-discovery records were not all individually reviewed. This is agent review, not human verification: no human recruiter, employer, or independent annotator confirmed the fields. Ambiguities remain unknown or unclear.

Bias and missingness

The earlier career/26fall dataset was used only for discovery and was biased toward Bay Area LLM roles; no applicant notes were imported. Fresh official sources broaden the sample but still favor accessible ATS pages. Salary and sponsorship disclosure are selective. A missing graduation date or experience statement stays unknown and never becomes zero. Seniority labels differ across employers; more senior postings do not imply more graduate openings.

Privacy and next improvement

The optional remember-on-device control stores a profile in the local browser only; there is no account or server-side applicant record in this product. Repeat the snapshot, record status changes, add sampled human adjudication, and compare extraction decisions against a written log.