

Methodology

AI Career Compass | visualizer rules and evidence boundaries

Snapshot	2026-09-21
Postings / companies	852 / 107
Source checks	198

Purpose and guided flow

AI Career Compass supports two views: Opportunity landscape describes the filtered sample and Skill lab compares local skill mentions. A lightweight tour prompts the student to explore opportunities, map skills, and try one suggested skill. Data & method keeps source links and course materials available.

Canonical snapshot

The snapshot contains 852 unique postings from 107 companies and 198 company/source checks. Sources are Ashby (309), Company careers (17), Greenhouse (522), Lever (4). Counts are postings, not headcount; source links, review labels, and data-quality notes remain available in the Data & method downloads.

Shared filters and measures

The two views share primary role direction, experience level, US region/state/remote, and graduate evidence. Landscape reports distributions and a skill-demand matrix. Skill lab reports the count and fraction of filtered postings mentioning at least one selected skill by role, then previews before/after coverage when one skill is tried.

Experience labels

Experience level is descriptive and independent from graduate eligibility. Existing title and curated labels remain. For records without those labels, cached official requirements map 0–2 years to entry, 3–4 to mid, 5+ to senior, prior professional experience without years to experienced, and explicit multiple seniority signals to open-level; unsupported cases stay unknown. The basis and explanation are recorded, while graduate flags and qualification paths are unchanged.

Skill interpretation

A posting counts once when it mentions at least one selected skill. Recommendations exclude selected skills and rank by posting mentions, then company breadth. Mention coverage is not proficiency, eligibility, readiness, hiring probability, or absence of an unmentioned skill. Automated extraction uses a fixed vocabulary.

Evidence and privacy

Records retain official URLs, source IDs, timestamps, extraction notes, and review level. The sample is non-representative and state/skill counts can overlap. Selected skills stay in page memory and optional device-local storage; they are not sent to a server. Public CSV downloads retain the evidence fields.