

Methodology

AI Career Compass | visualizer rules and evidence boundaries

Snapshot	2026-09-21
Postings / companies	852 / 107
Source checks	198

Purpose and interface

AI Career Compass helps students explore US AI career directions and choose skills to learn. Opportunity landscape shows the distribution of a dated posting sample. Skill lab compares selected skills with posting mentions and previews the effect of adding one skill. Data & method contains sources and course materials.

Canonical snapshot

The snapshot contains 852 unique postings from 107 companies and 198 company/source checks. Sources are Ashby (309), Company careers (17), Greenhouse (522), Lever (4). Counts are postings, not headcount; source links, review labels, and data-quality notes remain available in the Data & method downloads.

Global controls and views

The two main views share primary role direction, experience level (all seven levels or one level), US region/state/remote, and graduate evidence (all, explicit graduate, or early career). Landscape reports distributions by direction, experience, and region, plus a skill-demand matrix. Skill lab uses local-only skill chips, optionally remembered on this device, and reports the fraction and count of filtered postings mentioning at least one selected skill by role. Try one skill compares before and after posting coverage.

Skill interpretation

Skill lab measures mention coverage: a posting counts once when it mentions at least one selected skill. Missing skills are recommendations, ranked first by posting mentions and then company breadth. The top three are preparation prompts, followed by a small suggested learning project template. Mention coverage is not skill proficiency, eligibility, hiring probability, readiness, or a claim that an unmentioned skill is absent. Automated skill extraction uses a fixed vocabulary, so tracked terms are overrepresented.

Collection, provenance, and limits

Records come from official ATS or company-career pages and retain canonical URLs, source IDs, verification timestamps, extraction notes, and review level. The sample is non-representative: public ATS access, employer selection, posting turnover, and disclosure wording create coverage bias. Unknown skill is not skill absence; title level is not graduate eligibility; state and skill counts may overlap; salary is advertised base pay when disclosed. Earlier career/26fall material was discovery only and Bay Area LLM-biased; no applicant notes were imported.

Privacy and reproduction

Selected skills are held in page memory and, only when the user opts in, local browser storage; skill selection is not sent to a server or placed in a shared URL. Public CSV downloads retain the evidence fields. The repository includes collection, deduplication, analysis, and document-generation scripts for reproducing the snapshot.